




NATURAL LANGUAGE PROCESSING

Beller Stefan, Binder Lukas, Karasinski Maciej
Jacek, Kasper Bianca, Wichser Ilona



Outline



Dataset



LSTM



Dropout and LayerNormalization



GPUs



Models 1 - 9



Final network



Tracking of experiments

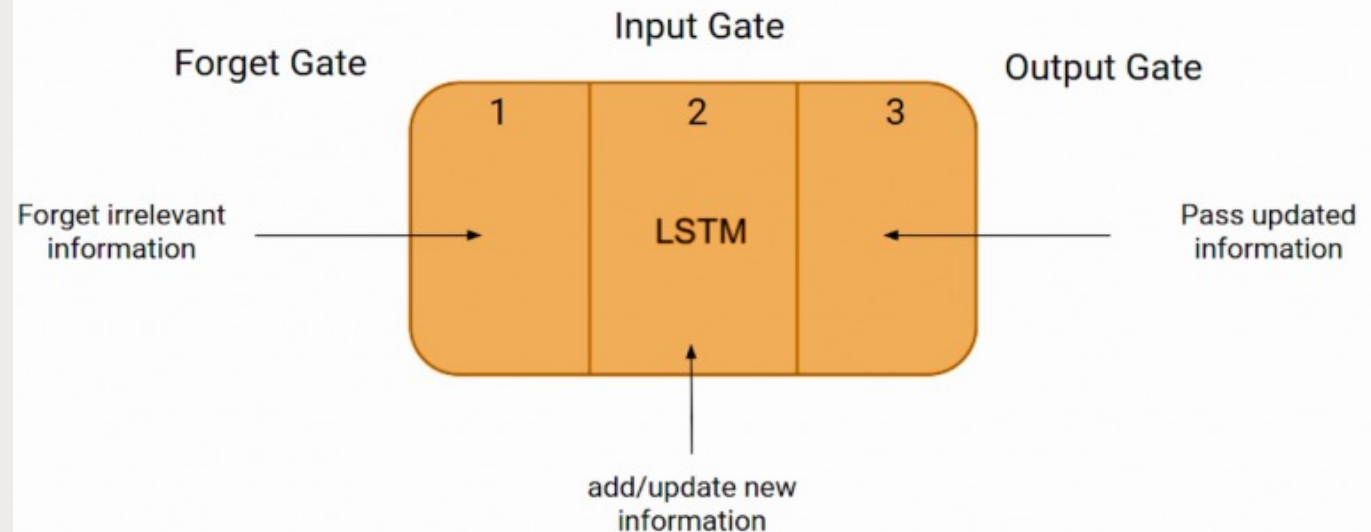
IMDB dataset for sentiment analysis

- Contains 50,000 movie reviews
- “good” and “negative” sentiment class labels
- Vocabulary size for our model: 5000 words

LSTM

- Used for all our models
- Special kind of RNN: capable of learning long-term dependencies
- Forgets, remembers and updates information

- 3 main gates



Why LSTM?



- Explicitly designed to avoid long-term dependency problems
- Useful for sentiment analysis of longer/ multiple sentences

Parameters:

- units: Dimensionality of output space
- return_sequences=True
 - *Enables output of previous LSTM layer to be used as an input to next LSTM layer*


Dropout vs. LayerNormalization

- Both to avoid overfitting
- **LayerNormalization:** Normalize the activations of the previous layer for each given example in a batch independently
- **Dropout:** Randomly sets input units to 0




GPUs

- Google Colab: Tesla T4



- Mac: M1 Pro



First networks



- Original: 1x model.add(SimpleRNN(100))

Model 1:

- 1x model.add(LSTM(100, return_sequences=True))
→ ~90% val_accuracy, ~50% test_accuracy

Model 2:

- 2x model.add(LSTM(100, return_sequences=True))
→ ~62% val_accuracy, 50% train_accuracy



Model 3

6m/epoch

embeddingSize = 64


- `model.add(LSTM(64, return_sequences=True))`
 - `model.add(Dropout(0.3))`
 - `model.add(LSTM(32, return_sequences=True))`
 - `model.add(Dropout(0.3))`
 - `model.add(LSTM(16, return_sequences=True))`
 - `model.add(Dropout(0.3))`
 - `model.add(LSTM(8, return_sequences=False))`
 - `model.add(LayerNormalization())`
-
- Test accuracy: 0.5009

Model 4

5-10m/epoch

embeddingSize = 128

- `model.add(LSTM(64, return_sequences=True, dropout=0.1))`
- `model.add(LayerNormalization())`
- `model.add(LSTM(32, return_sequences=True, dropout=0.1))`
- `model.add(LayerNormalization())`
- `model.add(LSTM(64, return_sequences=False))`
- `model.add(LayerNormalization())`
- OVERFITTING
- Test accuracy: 0.7838




Model 5 (lstm_3)

5-10m/epoch

embeddingSize = 128

- `model.add(LSTM(64, return_sequences=True, dropout=0.2))`
- `model.add(LayerNormalization())`
- `model.add(LSTM(32, return_sequences=True, dropout=0.2))`
- `model.add(LayerNormalization())`
- `model.add(LSTM(64, return_sequences=False))`
- `model.add(LayerNormalization())`
- OVERFITTING
- Test accuracy: 0.8174





Model 6 (lstm_2)

27-33s/epoch

embeddingSize = 128

- `model.add(LSTM(64, return_sequences=True))`
- `model.add(Dropout(0.2))`
- `model.add(LSTM(64, return_sequences=False))`
- `model.add(LayerNormalization())`
- Test accuracy: 0.8336
- Train accuracy: 0.9588
- Validation accuracy: 0.9219





Model 7 (lstm_do)

27-31s/epoch

embeddingSize = 128

- `model.add(LSTM(64, return_sequences=True))`
- `model.add(Dropout(0.2))`
- `model.add(LSTM(64, return_sequences=False))`
- `model.add(Dropout(0.2))`
- Test accuracy: 0.8722
- Train accuracy: 0.9433
- Validation accuracy: 0.9062

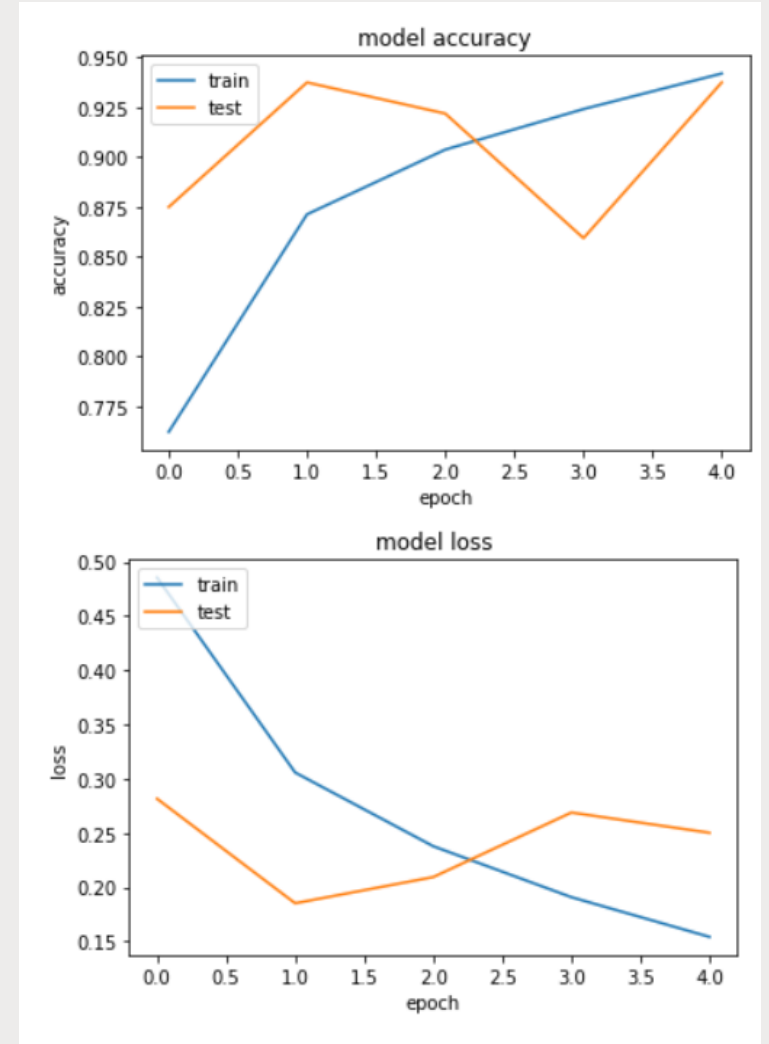


Model 8 (lstm_norm)

29-74s/epoch

embeddingSize = 128

- `model.add(LSTM(64, return_sequences=True))`
- `model.add(LayerNormalization())`
- `model.add(LSTM(64, return_sequences=False))`
- `model.add(LayerNormalization())`
- Test accuracy: 0.8687
- Train accuracy: 0.9420
- Validation accuracy: 0.9375




Model 9 (lstm_norm_256)

30-37s/epoch

embeddingSize = 256

- `model.add(LSTM(64, return_sequences=True))`
- `model.add(LayerNormalization())`
- `model.add(LSTM(64, return_sequences=False))`
- `model.add(LayerNormalization())`
- Test accuracy: 0.8330
- Train accuracy: 0.9661
- Validation accuracy: 0.9075



Final network: Model 7

- 2x LSTM layers, 2x Dropout
- Total parameters: 722,497
- Why selected?
 - *Highest test accuracy*
 - *Less overfitting than most other models*
- Embedding size = 128, because of community and papers suggestions
- Dropout, because better results

Deployment and tracking of experiments

- How has an experiment performed?
- We solved this by
 - Creating a branch for each experiment
 - Tracking the outcome of the branch
 - If an outcome is better than the main branch it is the new main branch

→ This was our simple way of tracking experiments



THANK YOU FOR YOUR ATTENTION.

QUESTIONS?

